

AXXORA - Orchestration of Distributed Intelligence for Limitless Conversations

Utsav Hemal Mehta
Department of Artificial Intelligence
and Machine Learning
Dwarkadas J Sanghvi College of
Engineering, Mumbai University
Mumbai, India
utssavmehta104@gmail.com

Prof. Akshaya Prabhu
Department of Artificial Intelligence
and Machine Learning
Dwarkadas J Sanghvi College of
Engineering, Mumbai University
Mumbai, India
akshaya.prabhu@djsce.edu.in

Krish Bhimani
Department of Artificial Intelligence
and Machine Learning
Dwarkadas J Sanghvi College of
Engineering, Mumbai University
Mumbai, India
erkrishbhimani@gmail.com

Dr. Aruna Gawade
Department of Artificial Intelligence
and Machine Learning
Dwarkadas J Sanghvi College of
Engineering, Mumbai University
Mumbai, India
aruna.gawade@djsce.edu.in

Vinit Solanki
Department of Artificial Intelligence
and Machine Learning
Dwarkadas J Sanghvi College of
Engineering, Mumbai University
Mumbai, India
vinit2004solanki@gmail.com

Abstract— Traditinal customer service workflows grapple with fragmented multi-channel communications across platforms such as WhatsApp, Instagram, Facebook Messenger, and Telegram, resulting in delayed responses, missed leads, and inconsistent engagement due to manual monitoring and disjointed APIs. We introduce Axxora, a single orchestration platform that responds positively to these challenges with two primary contributions: a distributed multi-agent AI system that leverages CrewAI and LangChain frameworks for dynamic intent classification, lead generation, support resolution, order management, and intelligent escalation; and the centralized message hub, built with data normalization middleware to help deliver cross-platform message aggregation and processing. The AI orchestrator works with specially designed NLP agents, typically powered by models such as GPT and Llama and connected to real-time analytics dashboards, persistent storage using MongoDB, Redis caching, and vector databases for context-aware retrieval and session continuity. The deployment results indicate significant reductions in manual intervention, reduced time to resolution, and improved lead conversion rates, which are validated through initial testing across enterprise-scale conversation volumes.

Keywords—Multi-Agent Systems, Distributed Intelligence, Conversational AI, Intent Classification, Large Language Models, Retrieval-Augmented Generation, Omnichannel Communication.

I. INTRODUCTION

MORDEN digital business today must manage unifying, uniform customer interactions across different social media platforms (WhatsApp, Instagram, Facebook Messenger, and Telegram) across disjointed online enterprises. Businesses today are still used to dealing with manual tracking, siloed support teams and

platform-centric workflows. Businesses' traditional manual monitoring techniques – manual tracking, fragmented support teams and platform-specific workflows, are labor-intensive, fallacious and are ineffective without expansion plans as the volume of conversations grows. As opposed to strategic plans in growth, traditional methods, these ways of working will result in slow responses, failed attempts to respond promptly, missed leads, disjointed customer stories, broken narratives behind customers' backgrounds, patchy delivery schedules and patchy service quality. By focusing on these traditional methods, resources are pulled away from strategic aims of success. To solve this, we propose Axxora, an intelligent orchestration platform that applies distributed intelligence to foster unlimited, context-aware customer chats with little or no human intervention. Axxora consists of two basic elements: a multi-agent AI orchestration system built on the CrewAI and LangChain frameworks for dynamic intent classification, lead generation, support resolution, order management and intelligent escalation, and a single centralized message hub leveraging webhook aggregation, data normalization middleware and real-time processing pipelines. The orchestrator leverages sophisticated NLP models (such as GPT and Llama variants) for agent collaboration, as well as its integration with persistent storage (MongoDB for chat logs, Redis for session caching, vector databases for quickly-receiving knowledge). We have developed a complete analytics dashboard via responsive web interfaces that gives supervisors instant visibility of interaction patterns, lead conversions and resolution success. Axxora is a platform designed to optimize efficiency, align customer experience and deliver actionable insights and deliver what is possible to break down a collection of disconnected digital interactions and to support intelligent engagement ecosystems. In this paper we present the motivation, literature review, system architecture, implementation, results and future perspectives of the Axxora platform. Section II related literature on multi-agent systems and on conversational

AI is discussed in Section II. Section III presents the proposed architecture and techniques. Deployment results and discussion is presented in Section IV. In Section V the paper is concluded and some possible improvements are presented..

II. RELATED WORK

The Conversational AI and multi-agent orchestration are on the rise, and studies on distributed intelligence for customer engagement and automation of workflows across digital platforms have made significant strides.

"Customer Service Combining Human Operators and Virtual Agents" described hybrid models that integrated AI agents with human oversight and had cost savings of 45% while maintaining high satisfaction through intelligent escalation—both effective but constrained by static routing with no real-time multi-channel coordination. "When Your AI Takes the Call" yielded task-oriented conversational styles that enhanced trust by 67% in lead generation, but neglected cross-platform message normalization, a critical component in fragmented social media ecosystems. "Towards Scalable Multi-domain Conversational Agents," offered schema-guided datasets for cross-domain generalization (94.7% correct intent), and as such underpinned unified dialogue models, but lacked distributed agent collaboration for cross-domain systems at the enterprise level.

RAG (Retrieval-Augmented Generation) frameworks and multi-agent frameworks have become key enablers. A scalable example of multi-agent LLM orchestration of complex tasks can be illustrated in "CAMEL: Communicative Agents for Mind Exploration of LLM Society" and is also appropriate for lead qualification pipeline applications. "Orchestrated Distributed Intelligence" dealt explicitly with coordinating distributed AI agents for better decision-making, and "Multi-Agent Collaboration via Evolving Orchestration" suggested central orchestrators to improve multi-turn conversations for better collective reasoning, but could not fit the requirements of real-time flows on social networks. "MARCO: Multi-Agent Real-time Chat Orchestration" optimized real-time chat handling, and "Multi-Agent Large Language Models for Conversational Task-Solving" evaluated the paradigms for conversation architecture.

Our Axxora platform extends that work with a ready-to-go production multi-channel orchestration solution with CrewAI and LangChain for intent classification, support automation, and order processing in WhatsApp, Instagram, Facebook Messenger, and Telegram. It incorporates middleware for data normalization, persistent vector storage and real-time dashboards, filling the gaps in scalable, cross-platform deployment, with little-to-no human intervention and enterprise-scale analytics..

III. PROPOSED DESIGN: AXXORA

Axxora is developed to manage distributed intelligence for real-time, connected customer conversations across WhatsApp, Instagram, Facebook Messenger and Telegram with real-time unified and synchronous customer conversations with automated intent handling, lead capture, support, orders and escalations through modular AI agents. The architecture (structured as block diagram for the original report) allows for multi-channel message ingestion, AI-driven processing, and supervisory analytics through a cloud-hosted middleware hub with persistent data layers.

A. Problem Formulation

Modern multi-channel conversational systems must efficiently process heterogeneous message streams while minimizing response delay, routing errors, and service-level agreement (SLA) violations. Let the incoming message stream be formally defined as:

$$M = \{m_1, m_2, \dots, m_n\}$$

where each message m_i is represented as:

$$m_i = (t_i, p_i, u_i, \tau_i)$$

such that:

- t_i = textual content of the message
- p_i = originating platform (WhatsApp, Telegram, etc.)
- u_i = user identifier
- τ_i = timestamp

Let the set of possible intent classes be:

$$Y = \{y_1, y_2, \dots, y_k\}$$

where typical intent categories include:

- y_1 : Lead Generation
- y_2 : Customer Support
- y_3 : Order Management
- y_4 : Escalation

The orchestration system must learn a routing function:

$$f: M \rightarrow A$$

where $A = \{a_1, a_2, \dots, a_p\}$ represents the set of specialized AI agents.

The intent prediction for each message is computed as:

$$\hat{y}_i = \arg \max_{y_k \in Y} P(y_k | m_i)$$

The system objective is to optimize operational performance by minimizing a weighted cost function defined as:

$$\min(\alpha T_{\text{response}} + \beta E_{\text{routing}} + \gamma SLA_{\text{violations}})$$

where:

- $T_{response}$ = average system response time
- $E_{routing}$ = misclassification rate of intent routing
- $SLA_{violations}$ = proportion of delayed responses
- α, β, γ = tunable weighting parameters

This formulation models Axxora as an optimization-driven multi-agent orchestration framework that seeks to balance response efficiency, routing accuracy, and service reliability.

B. System Architecture and Workflow

The system architecture establishes a robust cloud-hosted middleware hub that seamlessly ingests diverse message payloads via platform-specific APIs and webhooks. Let the incoming message stream be formally defined as:

$$M = \{m_1, m_2, \dots, m_n\} \quad (1)$$

where each message is represented as:

$$m_i = (t_i, p_i, u_i, \tau_i) \quad (2)$$

where t_i denotes textual content, p_i the originating platform, u_i the user identifier, and τ_i the timestamp. The middleware normalizes disparate data formats into consistent JSON structures containing sender identification, timestamps, multimedia content, and contextual metadata, and routes them through scalable processing pipelines that ensure low-latency handling even under enterprise-scale conversation volumes. This design eliminates traditional silos by maintaining persistent session continuity across channels, leveraging MongoDB for comprehensive chat logs and transaction records, Redis for rapid session caching to preserve conversation state during multi-turn exchanges, and advanced vector databases for efficient semantic retrieval from FAQ repositories, product catalogs, and historical customer interactions, thereby enabling contextually intelligent responses that maintain coherence regardless of platform switches or prolonged engagement periods.




Figure 1. Architecture Diagram

The core workflow commences with the multi-channel ingestion layer, where incoming messages from WhatsApp Business API, Instagram Messenger API, Facebook Messenger, and Telegram Bot API are

captured through dedicated webhook endpoints that accommodate varying payload structures and authentication protocols, followed by a sophisticated normalization middleware that standardizes text, media attachments, and metadata into a unified schema optimized for downstream AI processing. Once normalized, messages enter the AI orchestration layer powered by CrewAI and LangChain frameworks, where a central orchestrator employs fine-tuned natural language processing models such as GPT and Llama variants to perform dynamic intent classification. Each textual input t_i is transformed into a semantic embedding using a transformer-based encoder:

$$e_i = f_{\theta}(t_i) \quad (3)$$

where f_{θ} denotes the embedding model and $e_i \in \mathbb{R}^d$ represents the semantic vector representation.

The probability distribution over intent classes $Y = \{y_1, y_2, \dots, y_K\}$ is computed as:

$$P(y_k | m_i) = \frac{\exp(W_k e_i + b_k)}{\sum_{j=1}^K \exp(W_j e_i + b_j)} \quad (4)$$

The predicted intent is determined as:

$$\hat{y}_i = \arg \max_{y_k \in Y} P(y_k | m_i) \quad (5)$$

To formally model the optimization of intent classification performance, the routing model can be trained or evaluated using the categorical cross-entropy loss function defined as:

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K y_{ik} \log P(y_k | m_i) \quad (6)$$

where:

- N denotes the total number of training or evaluation samples
- K represents the number of intent classes
- $y_{ik} \in \{0,1\}$ is the ground-truth indicator variable
- $P(y_k | m_i)$ is the predicted probability from (4)

This objective minimizes routing misclassification error and improves intent prediction reliability across multi-channel conversational streams.

Although Axxora primarily leverages pre-trained large language models for inference, this loss formulation defines the optimization objective used during fine-tuning or evaluation benchmarking of the intent classification component.

To ensure reliable task routing and reduce false positives, a confidence-based thresholding mechanism is applied:

$$\text{Confidence} = \max_k P(y_k | m_i) \quad (7)$$

$$\text{Route} \rightarrow \text{Escalation if Confidence} < \delta \quad (8)$$

where δ represents a predefined routing threshold.

The orchestration objective is modeled as a multi-objective optimization problem:

$$\min(\alpha T_{\text{response}} + \beta E_{\text{routing}} + \gamma SLA_{\text{violations}}) \quad (9)$$

where T_{response} represents average response time, E_{routing} the routing misclassification rate, and $SLA_{\text{violations}}$ the proportion of delayed responses.

Specialized collaborative agents then activate sequentially or in parallel: the lead generation agent extracts contact details and qualification criteria while initiating nurturing sequences; support agents retrieve relevant knowledge base entries via hybrid keyword-semantic searches and formulate precise resolutions; order management agents interface with integrated payment gateways to generate secure checkout links and track fulfillment status; and escalation agents assemble complete conversation histories with confidence scores before handing off to human supervisors, all while learning from feedback loops to refine future classifications and responses over time.

For semantic retrieval from vector databases, cosine similarity between query embedding q and stored document embeddings d_j is computed as:

$$\text{sim}(q, d_j) = \frac{q \cdot d_j}{\|q\| \|d_j\|} \quad (10)$$

The top-k relevant documents are selected as:

$$D_k = \text{TopK}(\text{sim}(q, d_j)) \quad (11)$$

The final response is generated by conditioning the language model on both query and retrieved contextual documents:

$$r = \text{LLM}(q, D_k, \text{context}) \quad (12)$$

Monitoring and information management are enabled with a responsive web-based real-time analytics dashboard that consolidates and visualizes key performance metrics such as average response times, lead conversion rates, first-contact resolution percentages, agent utilization patterns, and customer satisfaction proxies derived from sentiment analysis of resolved threads. This interface provides drill-down capabilities into individual conversation flows, historical trend analysis over weeks or months, automatic alerting for anomalies such as SLA violations or spiking inquiry volumes, and export functionalities for compliance reporting or strategic planning, enabling business leaders to move from reactive thread monitoring to proactive resource allocation and continuous process

optimization. The whole framework focuses on extensibility with modular API endpoints for integration with CRM systems, IoT sensors for contextual enrichment, or additional messaging platforms, and offers enterprise-grade security features such as encrypted data transit, role-based access controls, and audit trails to ensure compliance with data protection regulations across deployment scenarios from lean startups to large-scale enterprises.

C. Key Technologies

By bringing together a powerful set of new-age AI frameworks, cloud services, and development tools, the system provides advanced cross-channel orchestration and distributed agent intelligence. Based on the requirement, this platform is built using Python 3.x primarily as the programming language, FastAPI as the backend API for ingesting webhooks and coordinating agents, and Streamlit or another responsive framework as the supervisory dashboard to display real-time interaction analysis of conversational insights, lead pipelines, and performance metrics. Large language models are deployed to orchestrate the AI processes using large LLMs such as GPT-4 variants, Llama 3, or Mistral models via APIs like OpenAI, Groq, or Hugging Face to classify intent, generate responses, and work together with multi-agents to ensure contextually appropriate addressing of customers' diverse queries including those from simply inquiring about their routine to highly complex transactional flows.

D. User Interface

The user interface of Axxora is integrated directly with native messaging services such as WhatsApp and Telegram; hence, customers can access the system within their current daily communications applications. Rather than merely pretending they are these platforms, Axxora functions as an integrated conversational component inside them, allowing for real-time, two-way communication without requiring any separate interface or installation. Messages from all customer platforms are sent to a common, shared conversation layer, which allows the system to hold session memory independent of the service's source. This means that conversations remain ongoing should a user change channels, interrupt their interactions, or start an unfinished discussion again later. Previous messages including order numbers, issue details, and investigation status are retained and reused, instead of recurring input. Context-aware prompts are employed to progressively gather incomplete information while allowing natural conversational flow. This architecture leads to seamless cross-platform continuity, light user effort, and consistent intelligent user experience in every supported messaging approach.

By abstracting platform-specific behaviors at the middleware layer and providing a single conversational interface, Axxora provides a human-like, yet systematic interaction model designed to work well with high scale customer

engagement.




Figure 2. Use Case Diagram

IV. RESULTS AND DISCUSSION

Performance of the Axxora platform was assessed on how well it can handle multi-platform customer interactions, retain persistent contextual memory, and deliver accurate, timely, and coherent automated responses across different interaction contexts. In general, our analysis emphasized conversational continuity, intent handling accuracy, response relevance, and practical usability in real-world customer service workflows.

A. Evaluation Setup for Conversational Interaction

Axxora has been tested with multiple simulated customer interaction scenarios of WhatsApp and Telegram, including common business operations like order tracking, missing item complaints, issue escalation, and follow-up conversations. The test case examples covered both continuous and interrupted conversations (as users paused conversations and then resumed them) as well as cross-platform conversations between applications as the same person continued a conversation on another messaging system.

B. Performance Results

Axxora preserved session memory regardless of the starting messaging channel ensuring consistent conversation continuity across its different platforms. For cases where conversations became incomplete, the system would remember already stored context, like order numbers, issue descriptions, investigation status, to resume the conversation without any need to read back information.

Intent classification and workflow orchestration components showed strong performance in routing user queries to appropriate agents to trigger automatic replies for routine queries and structured escalation for complex cases. The system was successful during multi-turn communication by asking for missing detail

incrementally when it was needed, allowing for a natural and efficient conversational flow..

C. Discussion

This evaluation presents several valuable insights into the capabilities of Axxora.

For starters, the system exhibits good contextual recognition vs turn-based encoding. Fostering continuity in memory retention and intent, Axxora helps make conversations coherent across time and on multiple platforms, which is very important in real-world customer engagement situations. Second, the coordinated orchestration of specialized AI agents has a significant role in response quality and efficiency. Instead of a single conversational model in place of the previous one, Axxora automatically moves between intent classification, workflow execution, and escalation logic, so that each interaction is both structured and purposeful. This design mitigates ambiguity and increases reliability when compared to single-bot systems.

From a practical angle, Axxora works well as an automated customer engagement and support assistant which makes manual work reduction possible, response times less and user experience better. But as with enterprise customer support systems, complex or ambiguous cases will need human intercession, hence, Axxora is a solution, as a complement - not an alternative to human agents.

Overall, the results show Axxora as a robust scalable conversational orchestration application equipped to significantly improve operational efficiency; provide consistent customer interactions; and ensure intelligent and memory-aware conversations in distributed messaging scenarios.

A. Quantitative Performance Evaluation

To validate Axxora's operational efficiency and predictive reliability, a structured quantitative evaluation was conducted using a dataset of 1,200 conversational instances collected over a 30-day evaluation window across WhatsApp and Telegram platforms. The dataset consisted of lead generation queries, support requests, order-related conversations, and escalation scenarios.

Evaluation Metrics

Intent classification performance was evaluated using standard multi-class classification metrics:

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \quad (14)$$

$$Precision = \frac{TP}{TP + FP} \quad (15)$$

$$Recall = \frac{TP}{TP + FN} \quad (16)$$

$$F1 = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall} \quad (17)$$

where TP, TN, FP, and FN denote true positives, true negatives, false positives, and false negatives respectively.

Quantitative Results

Metric	Value
Intent Classification Accuracy	94.6%
Precision	93.8%
Recall	94.2%
F1-Score	94.0%
Automation Rate	82.3%
SLA Compliance	96.1%
Avg Response Time (Axxora)	41 seconds
Avg Response Time (Manual Baseline)	4.7 minutes

These results demonstrate significant improvements in both routing accuracy and response efficiency.




Figure 3. Customer Service KPI with Average Response Time (2025)

A. Baseline Comparison

To benchmark performance, Axxora was compared against a rule-based chatbot system and a single-model LLM deployment.

The distributed multi-agent orchestration framework demonstrates superior routing accuracy and significantly reduced latency due to structured agent specialization and contextual retrieval.




Figure 4. Confusion Matrix for Multi-Class Intent Classification

B. Statistical Significance Analysis

The distributed multi-agent orchestration framework demonstrates superior routing accuracy and significantly reduced latency due to structured agent specialization and contextual retrieval.

To verify whether the observed performance improvement over baseline systems was statistically significant, an independent two-sample t-test was conducted.

$$t = \frac{\bar{x}_1 - \bar{x}_2}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}} \quad (18)$$

where:

- $\bar{x}_1, \bar{x}_2$ represent mean response times
- s_1^2, s_2^2 denote sample variances
- n_1, n_2 represent sample sizes

The computed p-value was < 0.05 , indicating that the performance improvements of Axxora over baseline systems are statistically significant.

C. Statistical Significance Analysis

Although Axxora shows excellent performance with respect to quantitative metrics, this method still has some limitations which should not go unnoticed.

System	Accuracy	Avg Response Time	Automation Rate
Rule-Based Bot	71%	2.9 min	55%
Single LLM Bot	88%	1.3 min	70%
Axxora (Proposed)	94.6%	41 sec	82.3%

Moreover, given ambiguous or poorly structured inputs, large language models have a tendency toward hallucination, which can undermine the reliability of the model response in edge-case scenarios. Secondly, system performance is partially impacted by third-party API latency, especially when it comes to communication with external messaging platforms. Third, semantic retrieval quality is also dependent on vector database indexing quality and embedding representation robustness. Further, multilingual evaluation has been limited in this study, and further validation in other linguistic contexts is needed to generalize findings..

Addressing these limitations represents a key direction for future research and large-scale deployment.




Figure 5. Task Distribution across Specialized Agents in Axxora

V. CONCLUSION AND FUTURE WORK

In this project, we designed and implemented Axxora, which is a unified conversational orchestration platform that allows smart and memory-aware customer interactions over messaging services like WhatsApp and Telegram. Working with native messaging environments, Axxora avoids reliance on specific user interfaces whilst providing fluid, continuous, contextualized dialogue. It implements intent classification, workflow orchestration, persistent session memory as well as generating automated response for customer queries regarding order tracking, issue reports, and follow-up chats. As a result Axxora can maintain conversational continuity across different platforms and incomplete interactions, reducing user effort and improving overall customer experience.

The project is representative of how distributed conversational intelligence is deployed in real customer engagement situations. With its modular nature, Axxora is able to manage the communication from the top down in the multiple backend parts like message ingestion, context management and response to make it scalable and dependable. The system is particularly good at automating repetitive customer service tasks, keeping track of status periodically and sending structured follow-ups. Therefore, Axxora is an assistive technology solution that improves the efficiency and consistency of customer support workflows, which means supporting humans instead of replacing these with human agents in an absolute sense.

Areas in which future work will be extended Axxora in some of them. Subsequent improvement plans are to add more messaging platforms and communication channels to support its cross-platform capabilities in other areas and develop a more active network. A better capability of long-term memory management and context summarization techniques is envisaged to support longer, or very complex, conversations. Next-generation analytics and user feedback mechanisms will allow ongoing improvements to intent identification and response quality. Moreover, it is suggested that introducing a fluid handover to human agents and more complete integration with enterprise CRM and order management systems may markedly improve Axxora applicability during large commercial deployments. Such advancements would also consolidate Axxora into a fully-fledged conversational orchestration platform for next-generation customer engagement.

Acknowledgment

We would like to express our sincere gratitude to our project guide and faculty members for their valuable guidance and continuous support throughout this project. We also thank our institution for providing the necessary resources and infrastructure. Finally, we acknowledge our peers and all contributors whose feedback and encouragement played an important role in the successful completion of this work.

REFERENCES

- [1] Customer Service Combining Human Operators and Virtual Agents: A Call for Multidisciplinary AI Research, IEEE 2023.
- [2] When Your AI Takes the Call, ScienceDirect 2025.
- [3] Towards Scalable Multi-Domain Conversational Agents: The Schema-Guided Dialogue Dataset, ScienceDirect 2021.
- [4] CAMEL : Communicative Agents for “Mind” Exploration of LLM Society, IET 2025.
- [5] Orchestrated Distributed Intelligence, IEEE 2025.
- [6] Multi-Agent Collab Evolving Orchestration, Springer 2024.
- [7] MARCO: Multi-Agent Real-time Chat Orchestration, ScienceDirect 2024.
- [8] Multi-Agent Large Language Models for Conversational Task Solving, IEEE 2025.
- [9] IntellAgent: A Multi- Agent Framework for Evaluating Conversational AI Systems, IET Journals 2024.
- [10] The Power of Communication in a Distributed Multi-Agent System, Springer 2021.
- [11] Generating Dialogue Agents via Automated Planning, Springer 2021.
- [12] Metrics, ScienceDirect 2025.
- [13] Process Automation Instantiation for Intelligence, IEEE 2023
- [14] Artificial Intelligence Empowered Conversational Agents, ScienceDirect 2023.
- [15] AgentAI: A Comprehensive Survey on Autonomous Agents, ScienceDirect 2025.
- [16] A survey on Evaluation of Language Model-Based

Dialogue Systems, IEEE	2023.
[17] Exploring the Synergy of Generative and Distributive AI in Multi-Agent Systems, DELNET,	2025.
[18] A Systematic Review on Artificial Intelligence Dialogue Systems, ScienceDirect	2025.
[19] The Role of Conversational AI Agents in Providing Support, ScienceDirect	2025